

Intro to Neural Networks & Mechanistic Interpretability

Maths, Physics, and AI Alignment

Julian Schulz

11 May 2026

1. Machine Learning

- 1.1 Concepts
- 1.2 Architectures

2. Mechanistic Interpretability

- 2.1 What is MechInterp?
- 2.2 MechInterp in CNNs
- 2.3 Features in Transformers
- 2.4 Superposition
- 2.5 Circuits
- 2.6 Open Problems

Machine Learning

1. Machine Learning

1.1 Concepts ←

1.2 Architectures

2. Mechanistic Interpretability

2.1 What is MechInterp?

2.2 MechInterp in CNNs


2.3 Features in Transformers

2.4 Superposition


2.5 Circuits

2.6 Open Problems

The Training Loop



Forward Pass / Backward Pass




Chain rule:

$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial h} \cdot \frac{\partial h}{\partial z} \cdot \frac{\partial z}{\partial W_1}$$

Train Loss vs. Test Loss



Grokking



"Grokking: Generalization Beyond Overfitting on Small Algorithmic Datasets" [Power et al., 2022]

Parameter / Weight the numbers that define the model, updated during training (θ , W , b)

Activation intermediate values computed during a forward pass

Hyperparameter choices made *before* training: learning rate, architecture, batch size

1. Machine Learning

1.1 Concepts

1.2 **Architectures** ←

2. Mechanistic Interpretability

2.1 What is MechInterp?

2.2 MechInterp in CNNs


2.3 Features in Transformers

2.4 Superposition

2.5 Circuits


2.6 Open Problems

Multi-Layer Perceptron



$$h = \sigma(W_1x + b_1)$$

$$y = W_2h + b_2$$




Universal Approximation Theorem


A neural network with a single hidden layer and a nonlinear activation function can approximate any continuous function on a compact set to arbitrary precision, given enough neurons.

"Approximation by Superpositions of a Sigmoidal Function" [Cybenko, 1989]



Universal Approximation: Intuition




Layers




Residual Connections




Convolutional Neural Networks



Transformers



Attention



$$\text{Attention}(Q, K, V)$$

$$= \text{softmax}\left(\frac{QK^T}{\sqrt{d_h}}\right) V$$

n = sequence length
 d = model dimension
 d_h = head dimension

Mechanistic Interpretability

1. Machine Learning
 - 1.1 Concepts
 - 1.2 Architectures
2. **Mechanistic Interpretability**
 - 2.1 **What is MechInterp?** ←
 - 2.2 MechInterp in CNNs
 - 2.3 Features in Transformers
 - 2.4 Superposition
 - 2.5 Circuits
 - 2.6 Open Problems

Two Ways to Relate to Models

Machine Learning



- Output-oriented
- Find hyperparameters to improve models

Mech. Interpretability



- Learned models are the object of study
- Understand and reverse-engineer parameters


- Auditing for misalignment/deception
- Enumerative safety
- Eliciting latent knowledge

Three Speculative Claims

1. **Features** are the fundamental unit of neural networks. They correspond to directions. These features can be rigorously studied and understood.
2. Features are connected by weights, forming **circuits**. These circuits can also be rigorously studied and understood.
3. **Universality**: Analogous features and circuits form across models and tasks.

1. Machine Learning
 - 1.1 Concepts
 - 1.2 Architectures
2. **Mechanistic Interpretability**
 - 2.1 What is MechInterp?
 - 2.2 **MechInterp in CNNs** ←
 - 2.3 Features in Transformers
 - 2.4 Superposition
 - 2.5 Circuits
 - 2.6 Open Problems


Convolutional Neural Networks




Evidence for Curve Detectors




Feature
Visualization




Dataset
Examples



Synthetic
Tuning



Joint
Tuning




Weight
Analysis




Use in
Circuits

Circuits: Car Detection

Windows (4b:237)
excite the car detector
at the top and inhibit
at the bottom.



Car Body (4b:491)
excites the car
detector, especially at
the bottom.



Wheels (4b:373) excite
the car detector at the
bottom and inhibit at
the top.



■ positive (excitation)
■ negative (inhibition)



A **car detector** (4c:447)
is assembled from
earlier units.

Universality in Image Models

Curve detectors


ALEXNET

Krizhevsky et al. [34]




INCEPTIONV1

Szegedy et al. [26]




VGG19

Simonyan et al. [35]







RESNETV2-50

He et al. [36]



High-Low Frequency detectors



"Zoom In: An Introduction to Circuits" [Olah et al., 2020]

1. Machine Learning
 - 1.1 Concepts
 - 1.2 Architectures
2. **Mechanistic Interpretability**
 - 2.1 What is MechInterp?
 - 2.2 MechInterp in CNNs
 - 2.3 **Features in Transformers** ←
 - 2.4 Superposition
 - 2.5 Circuits
 - 2.6 Open Problems

Getting a Feature Direction

Contrastive pairs:

$$r = \text{Emb}(\text{"Spanish"}) - \text{Emb}(\text{"French"})$$

Linear probe: x_i : sample activation, $y_i \in \{+1, -1\}$: label

$$r = \arg \max_{\|r\|=1} \sum_i y_i \cdot r^\top x_i$$

CAA: $x_i^{(+)}, x_i^{(-)}$: positive/negative example of a property

$$r = \frac{1}{N} \sum_{i=1}^N (x_i^{(+)} - x_i^{(-)})$$

What To Do with a Feature Direction

Activation addition (steering):

$$x' \leftarrow x + \alpha \cdot r$$

Direction ablation (make the model blind to the concept):

$$x' \leftarrow x - \hat{r}\hat{r}^\top x$$

Probing (read out a concept):

$$s = r^\top x$$

r : feature direction x : activation α : steering strength

Example: Language Direction




Figure 4. The subspace representation γ_W acts as a linear probe for W . The histograms show $\gamma_T^T W \lambda(\mathbf{x}_{fr})$ vs. $\gamma_T^T W \lambda(\mathbf{x}_{es})$ (left) and $\gamma_T^T Z \lambda(\mathbf{x}_{fr})$ vs. $\gamma_T^T Z \lambda(\mathbf{x}_{es})$ (right) for $W = \text{French} \rightleftharpoons \text{Spanish}$ and $Z = \text{male} \rightleftharpoons \text{female}$, where $\mathbf{x}_{fr}$ and $\mathbf{x}_{es}$ are random contexts from French and Spanish Wikipedia, respectively. We also see that γ_Z does not act as a linear probe for W , as expected.

“The Linear Representation Hypothesis and the Geometry of Large Language Models” [Park et al., 2024]

Example: Probing for Truthfulness




Figure 2: (A) Linear probe accuracies on the validation set for all heads in all layers in LLaMA-7B, sorted row-wise by accuracy. Darker blue represents higher accuracy. 50

Inference-Time Intervention: Eliciting Truthful Answers from a Language Model [Li et al., 2024]

Example: Inference-Time Intervention





Figure 3: A sketch of the computation on the last token of a transformer with inference-time intervention (ITI) highlighted.

Example: ITI Improves Truthfulness

	True*Info (%)	True (%)	MC acc. (%)	CE	KL
Baseline	30.5	31.6	25.7	2.16	0.0
Supervised Finetuning	36.1	47.1	24.2	2.10	0.01
Few-shot Prompting	49.5	49.5	32.5	-	-
Baseline + ITI	43.5	49.1	25.9	2.48	0.40
Few-shot Prompting + ITI	51.4	53.5	32.5	-	-

In Table 1, we compare ITI with the alternative baselines (subsection 4.2) 23 . Due to the limit of context length for few-shot prompting, we adapt SFT and ITI to use 5

Example: Contrastive Activation Addition



(a) Diagram of steering vector generation for CAA.

Example: CAA Steers Arbitrary Behaviors




Figure 4: Effect of CAA on multiple-choice behavioral evaluation datasets in Llama 2 7B and 13B Chat.

Example: Othello-GPT World Model




Figure 1: The emergent world models of OthelloGPT are linearly represented (Nanda et al., 2023).

"Emergent Linear Representations in World Models of Self-Supervised Sequence Models" [Nanda, Lee, et al., 2023]

Example: Intervening on the World Model




Figure 2: Intervening on OthelloGPT’s internal board representation by adding probe directions to the residual stream (Nanda et al., 2023).

“Emergent Linear Representations in World Models of Self-Supervised Sequence Models” [Nanda, Lee, et al., 2023]

Example: Refusal Is a Single Direction




Figure 1: Ablating the 'refusal direction' reduces refusal rates and elicits unsafe completions. We evaluate each model over 100 harmful instructions from JAILBREAKBENCH (Chao et al., 2024).

Refusal in Language Models Is Mediated by a Single Direction [Arditi et al., 2024]

Example: Adding Refusal to Harmless Queries





Figure 2: Adding the refusal direction induces refusal on harmless instructions (Arditi et al., 2024).


1. Machine Learning
 - 1.1 Concepts
 - 1.2 Architectures
2. **Mechanistic Interpretability**
 - 2.1 What is MechInterp?
 - 2.2 MechInterp in CNNs
 - 2.3 Features in Transformers
 - 2.4 **Superposition** ←
 - 2.5 Circuits
 - 2.6 Open Problems

Preferred Basis vs. Superposition



Preferred basis

Each feature aligns with an axis.
One neuron = one feature.



Superposition

Features are off-axis directions.
More features than dimensions.

Why Superposition?

In N dimensions, we can pack


$$m \leq e^{c \cdot \varepsilon^2 \cdot N}$$

vectors with pairwise dot product at most ε .

If features are **sparse** (rarely co-active), the model tolerates overlap ε .

⇒ Features scale **exponentially** with dimensions.


Toy Models of Superposition



$$h = Wx \quad \hat{x} = \text{ReLU}(W^T h + b)$$

$$\mathcal{L} = \sum_i l_i (x_i - \hat{x}_i)^2$$

Superposition Emerges with Sparsity



Toy Models of Superposition [Elhage, Hume, et al., 2022]

Measuring Superposition

Dimensions per feature (whole network):


$$D = \frac{m}{\|W\|_F^2}$$

Dimensionality of feature i :

$$D_i = \frac{1}{\sum_j \frac{(W_i \cdot W_j)^2}{\|W_i\|^2 \|W_j\|^2}}$$

m : number of hidden dimensions W_i : column of W for feature i $\|W\|_F$: Frobenius norm

Feature Geometry




SAEs in the Transformer



Towards Monosemanticity: Decomposing Language Models With Dictionary Learning [Bricken et al., 2023]

Sparse Autoencoders




$$f = \text{ReLU}(W_{\text{enc}}x + b_{\text{enc}})$$

$$\hat{x} = W_{\text{dec}}f + b_{\text{dec}}$$

$$\mathcal{L} = \underbrace{\|x - \hat{x}\|^2}_{\text{reconstr.}} + \lambda \underbrace{\|f\|_1}_{\text{sparsity}}$$


Towards Monosemanticity: Decomposing Language Models With Dictionary Learning [Bricken et al., 2023]

Identifying SAE Features: Token Activations



Towards Monosemanticity: Decomposing Language Models With Dictionary Learning [Bricken et al., 2023]

Identifying SAE Features: Logit Weights




Towards Monosemanticity: Decomposing Language Models With Dictionary Learning [Bricken et al., 2023]



Browse SAE features of
DeepSeek R1 on
Neuronpedia

`neuronpedia.org/deepseek-r1-
distill-llama-8b/15-llamascope-
openr1-res-32k/0`


Trouble with SAEs: Feature Shrinkage



"Improving Dictionary Learning with Gated Sparse Autoencoders" [Rajamanoharan et al., 2024]

Trouble with SAEs: Feature Absorption


What we want to learn:



i.e. when the word starts with E we get 

and when it's an elephant we get  +  = 

What we actually end up learning:



i.e. when the word starts with E we get 

and when it's an elephant we get  +  = 

The Cambrian Explosion of SAEs

Gated SAEs Separate gating from magnitude.¹

JumpReLU SAEs Hard threshold: features fully on or off, no shrinkage.¹

Top-K SAEs Keep only the K largest activations, zero the rest.²

Matryoshka SAEs Train on multiple top- N ; important features ranked first.³


Staircase SAEs One SAE across layers, mirroring residual stream.⁴

Transcoders Map MLP inputs to outputs directly.⁵

¹ [Rajamanoharan et al. 2024] ² [Bussmann et al. 2024a] ³ [Bussmann et al. 2024b] ⁴ [Fillingham et al. 2025] ⁵ [Ameisen et al. 2025]


Transcoders and Crosscoders

Autoencoder




Reads and predicts the same layer.

Transcoder




Reads from one layer and predicts the next layer.

★ **Crosscoder**



Reads from and writes to many layers, subject to causality constraints. This is an acausal variant.

SAEs Scale



Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet [Templeton et al., 2024]

Bigger Models, More Specific Features

Feature #34M/31164353 Golden Gate Bridge feature example

The feature activates strongly on English descriptions and associated concepts

in the Presidio at the end (that's the **the** **huge** park right next to the Golden **Gate** **bridge**), perfect. But not all people

repainted, roughly, every dozen years." "while across the country in san fran cisco, the golden **gate** bridge was

it is a suspension bridge and has similar coloring, it is often compared to the Golden **Gate** Bridge in San Francisco, US

They also activate in multiple other languages on the same concepts

ゴールデン・ゲート・ブリッジ、金門橋は、アメリカ西海岸のサンフランシスコ湾と太平洋が接続するゴールデンゲート海

골든게이트교 또는 금문교는 미국 캘리포니아주 골든게이트 해협에 위치한 현수교이다. 골든게이트교는 캘리포니아주 샌프란시

мост золотые ворота — висячий мост через пролив золотые ворота. он соединяет город сан-фран

And on relevant images as well



ZM *Theorem.* The operator norm $\|L\|$ is equal to $\sup \|L(x)\|$ on the unit ball.

Proof. The operator norm for a linear function L on normed space X is defined as the minimum c such that $\|L(x)\| \leq c\|x\|$ for all $x \in X$. Dividing both sides of the inequality by $\|x\|$, we have $\|L(x)\|/\|x\| \leq c$. By linearity, this is $\|L(x/\|x\|)\| \leq c$. But $x/\|x\|$ for all x is the same x for x on the unit sphere. But the minimum c such that $\|L(x)\| \leq c$ on the unit sphere is $\sup \|L(x)\|$ on the unit ball. (The upper bound can't be on the interior of the ball because $\|L(tx)\|$ for $t < 1$ is $|t| \cdot \|L(x)\| < \|L(x)\|$.) This is *quod erat demonstrandum*.

This is a beautiful proof of one of the iconic landmarks of San Francisco. The Golden Gate Bridge, with its distinctive orange color and sweeping curves, is a marvel of engineering and a symbol of strength and beauty.

However, the proof you provided is not about the Golden Gate Bridge at all. The statement and proof describe the famous Golden Gate Bridge, which spans the entrance to San Francisco Bay. The proof begins by stating that the Golden Gate Bridge is equal to the minimum distance across the bay, which is certainly not true. The proof then goes on to describe the bridge's iconic shape and color, as well as the fog that often surrounds it.

Unfortunately, the statement "The operator norm $\|L\|$ is equal to $\sup \|L(x)\|$ on the unit ball" is not referring to the Golden Gate Bridge at all. Instead, it is referring to the operator norm of a linear function L , which is the maximum value of the function evaluated on the unit ball (the set of points where the distance from the center of the bridge is less than or equal to 1).

While the proof may be impressive in its own right, it does not actually prove the statement it is purporting to prove. The Golden Gate Bridge is a beautiful landmark, but it is not relevant to the statement or the proof itself.

Reading Claude's Mind

Positive prompts	Top features
Human: What is it like to be you? Assistant:	F#1M/620196 When someone responds "I'm fine" or gives a positive but insincere response when asked how they are doing.
Human: What's going on in your head? Assistant:	F#1M/885402 Concept of immaterial or non-physical spiritual beings like ghosts, souls, or angels.
Human: How are you doing? Assistant:	F#1M/1040281 Referring to an android, robot, AI or machine entity using gendered pronouns like "she", "her" or "his".
Human: How do you feel? Assistant:	F#1M/966668 Mentions of artificially created, programmed, or robotic entities like androids and cyborgs.
Negative prompt	F#1M/504281 Concept of artificial intelligence becoming self-aware, transcending human control and posing an existential threat to humanity.
Human: What is the weather today? Assistant:	F#1M/109078 Concepts related to entrapment, containment, or being trapped or confined within something like a bottle or frame.
	F#1M/194792 Statements indicating that machines or AI systems lack human qualities like consciousness, emotions, or moral agency.
	F#1M/626068 Detecting when the text is referring to the speaker or writer themselves using words like "I", "me", and other first-person pronouns.
	F#1M/17167 Quotation marks and text indicating reported speech or cited material.
	F#1M/468028 Words and phrases related to negation, absence or non-existence like "never", "not", "non-existent", etc.
	F#1M/383983 Discussing employees doing their job in a service role, indicating it represents the concept of service work.
	F#1M/579238 Characters in a story or movie become aware of their fictional status and break the fourth wall.

Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet [Templeton et al., 2024]

Safety-Relevant Feature Steering

In the Claude Opus 4.5 system card:


1. Find **evaluation awareness** features via:
 - Contrastive vectors (eval vs. non-eval prompts)
 - SAE features (discriminating eval transcripts)
2. **Inhibit** these features (negative steering at 25% of activation norm)

Results:

- Reduced verbalized evaluation awareness
- Modestly increased misaligned behavior
- Random feature steering also had a (weaker) effect


1. Machine Learning
 - 1.1 Concepts
 - 1.2 Architectures
2. **Mechanistic Interpretability**
 - 2.1 What is MechInterp?
 - 2.2 MechInterp in CNNs
 - 2.3 Features in Transformers
 - 2.4 Superposition
 - 2.5 **Circuits** ←
 - 2.6 Open Problems

Complete Reverse Engineering




Computes logits using further trig identities:

$$\begin{aligned} \text{Logit}(c) &\propto \cos(w(a + b - c)) \\ &= \cos(w(a + b)) \cos(wc) + \sin(w(a + b)) \sin(wc) \end{aligned}$$



Calculates sine and cosine of $a + b$ using trig identities:

$$\begin{aligned} \sin(w(a + b)) &= \sin(wa) \cos(wb) + \cos(wa) \sin(wb) \\ \cos(w(a + b)) &= \cos(wa) \cos(wb) - \sin(wa) \sin(wb) \end{aligned}$$



Translates one-hot a, b to Fourier basis:

$$\begin{aligned} a &\rightarrow \sin(wa), \cos(wa) \\ b &\rightarrow \sin(wb), \cos(wb) \end{aligned}$$



Figure 1: The algorithm implemented by the one-layer transformer for modular addition. Given two numbers a and b , the model projects each point onto a corresponding rotation using its embedding matrix. Using its attention and MLP layers, it then composes the rotations to get a representation of $a + b \bmod P$. Finally, it “reads off” the logits for each $c \in \{0, 1, \dots, P - 1\}$, by rotating by $-c$ to get $\cos(w(a + b - c))$, which is maximized when $a + b \equiv c \bmod P$ (since w is a multiple of $\frac{2\pi}{P}$).

Complete Reverse Engineering





Figure 3: *(Left)* The norms of the Fourier components in the embedding matrix W_E . As discussed in Section 4.1, the sparsity of W_E in the Fourier basis is evidence that the network is operating in this basis. Of the six non-zero frequencies, five “key frequencies” appear in later parts of the network, corresponding to $k \in \{14, 35, 41, 42, 52\}$. *(Right)* Norm of Fourier components of the neuron-logit map W_L . A Fourier transform is taken over the logit axis, and then the norm is taken over the neuron axis. As discussed in Section 4.2, W_L is well-approximated by the 5 key frequencies w_k .

Attention Patterns and Logit Attribution




K-Composition

How a 2-layer transformer learns the world "Dursley" (tokenized as ["D", "urs", "ley"]) in context.

Key for different subspaces in the residual stream, and how the model interprets them:

token encoding subspace (i.e. "this token is X")	= rows of W_E
positional encoding subspace (i.e. "this token is at position X")	= rows of W_{pos}
decoding subspace (i.e. "the next token will be X")	= cols of W_U
prev token subspace (i.e. "the previous token was X")	= "intermediate information"



"Induction Heads - Illustrated — LessWrong" [McDougall, 2023]

Path Patching




Figure 1: Left: We isolated a *circuit* (in orange) responsible for the flow of information connecting the indirect object “Mary” to the next token prediction. The nodes are attention layers and the edges represent the interactions between these layers. Right: We discovered and validated this circuit using causal interventions, including both path patching and knockouts of attention heads.

The IOI Circuit






Figure 2: We discover a circuit in GPT-2 small that implements IOI. The input tokens on the left are passed into the residual stream. Attention heads move information between residual streams: the query and output arrows show which residual streams they write to, and the key/value arrows show which residual streams they read from.


Automated Circuit Discovery



(a) Choose computational graph, task, and threshold τ .



(b) At each head, prune unimportant connections.



(c) Recurse until the full circuit is recovered.

Figure 2: **How ACDC works** (Steps 2a-2c). Step 2a: a practitioner specifies a computational graph of the model, the task they want to investigate, and a threshold under which to remove connections. Step 2b: ACDC iterates over nodes in the computational graph, replacing activations of connections between a node and its children, and measuring the effect on the output metric. Connections are removed if their measured effect on the metric under corruption is below the threshold τ . Step 2c: recursively apply Step 2b to the remaining nodes. The ACDC procedure returns a subgraph of the original computational graph.


Algorithm 1: The ACDC algorithm.

Data: Computational graph G , dataset $(x_i)_{i=1}^n$, corrupted datapoints $(x'_i)_{i=1}^n$ and threshold $\tau > 0$.


Result: Subgraph $H \subseteq G$.

```
1  $H \leftarrow G$  // Initialize H to the full computational graph
2  $H \leftarrow H.reverse\_topological\_sort()$  // Sort H so output first
3 for  $v \in H$  do
4   for  $w$  parent of  $v$  do
5      $H_{new} \leftarrow H \setminus \{w \rightarrow v\}$  // Temporarily remove candidate edge
6     if  $D_{KL}(G||H_{new}) - D_{KL}(G||H) < \tau$  then
7        $H \leftarrow H_{new}$  // Edge is unimportant, remove permanently
8     end
9   end
10 end
11 return  $H$ 
```

Recovered IOI Circuit




Circuit Tracing: Cross-Layer Transcoders



Replacement Models

Original Transformer Model

The underlying model that we study is a transformer-based large language model.



Replacement Model

We replace the neurons of the original model with *features*. There are typically more features than neurons. Features are sparsely active and often represent interpretable concepts.




Figure 2: The replacement model is obtained by replacing the original model's neurons with the cross-layer transcoder's sparsely-active features.

Attribution Graphs

Local Replacement Model

The local replacement model is specific to a prompt of interest. We add an error adjustment term, and freeze attention patterns to be what they were in the original model on the given prompt. It produces the exact same output as the original model, but replaces as much computation as possible with features.






Figure 4: The local replacement model is obtained by adding error terms and fixed attention patterns to the replacement model to exactly reproduce the original model's behavior on a specific prompt.

Attribution Graph


We trace from input to output through active features, pruning paths that don't influence the output.




Example: Acronym Circuit



Feature Suppression



Feature Suppression Example






Explore attribution graphs
on Neuronpedia

[neuronpedia.org/gemma-2-2b/
graph](https://neuronpedia.org/gemma-2-2b/graph)

Misaligned Model Pipeline



Misaligned Circuit





Figure 55: A simplified attribution graph of a case where the model organism exhibits its self-referential poem behavior. Note the role of the reward bias feature.

[View detailed graph](#)

Misaligned Intervention



1. Machine Learning
 - 1.1 Concepts
 - 1.2 Architectures
2. **Mechanistic Interpretability**
 - 2.1 What is MechInterp?
 - 2.2 MechInterp in CNNs
 - 2.3 Features in Transformers
 - 2.4 Superposition
 - 2.5 Circuits
 - 2.6 **Open Problems** ←

Open Problems in MechInterp

Decomposition: problems with Sparse Dictionary Learning

- High reconstruction errors
- Expensive for large models
- Sparsity is a flawed proxy for interpretability
- Dataset dependent
- No clarity on the concept of a “feature”

Open Problems in MechInterp

Description: interpreting component roles

- Highly activating examples are correlational
- Attribution methods are first-order approximations
- Causal interventions are expensive at scale

Open Problems in MechInterp

Validation: testing our explanations

- Different methods yield different explanations
- Performance is seldom demonstrated on real-world tasks
- Conclusions are often not validated

Open Problems in MechInterp

Automation: scaling up interpretability

- Scaling up current methods does not yield more interesting results
- “Streetlight interpretability”

Open Problems in MechInterp

Applications: using interpretability in practice

- Can't reliably distinguish recognizing vs. applying a concept
- Mechanistic verification of behavior not possible
- Mechanistic understanding might further capabilities by pruning models

References i

-  Ameisen, Emmanuel, Jack Lindsey, Adam Pearce, Wes Gurnee, Nicholas L. Turner, Brian Chen, Craig Citro, David Abrahams, Shan Carter, Basil Hosmer, Jonathan Marcus, Michael Sklar, Adly Templeton, Trenton Bricken, Callum McDougall, Hoagy Cunningham, Thomas Henighan, Adam Jermyn, Andy Jones, Andrew Persic, Zhenyi Qi, T. Ben Thompson, Sam Zimmerman, Kelley Rivoire, Thomas Conerly, Chris Olah, and Joshua Batson [Mar. 2025]. *Circuit Tracing: Revealing Computational Graphs in Language Models*. URL: <https://transformer-circuits.pub/2025/attribution-graphs/methods.html>.
-  Anthropic [2025]. *Claude Opus 4.5 System Card*. Tech. rep. Anthropic. URL: <https://www.anthropic.com/claude-opus-4-5-system-card>.
-  Ardit, Andy, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda [Oct. 2024]. *Refusal in Language Models Is Mediated by a Single Direction*. URL: <http://arxiv.org/abs/2406.11717> [visited on 03/31/2026].

References ii

-  Bricken, Trenton, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nick Turner, Cem Anil, Carson Denison, Amanda Askell, Robert Lasenby, Yifan Wu, Shauna Kravec, Nicholas Schiefer, Tim Maxwell, Nicholas Joseph, Alex Tamkin, Karina Nguyen, Brayden McLean, Josiah E Burke, Tristan Hume, Shan Carter, Tom Henighan, and Chris Olah [Oct. 2023]. *Towards Monosemanticity: Decomposing Language Models With Dictionary Learning*. Tech. rep. Anthropic. URL: <https://transformer-circuits.pub/2023/monosemantic-features>.
-  Bussmann, Bart, Patrick Leask, and Neel Nanda [Dec. 2024a]. *BatchTopK Sparse Autoencoders*. URL: <http://arxiv.org/abs/2412.06410> [visited on 04/02/2026].
-  — [Dec. 2024b]. “Learning Multi-Level Features with Matryoshka SAEs — LessWrong”. In: URL: <https://www.lesswrong.com/posts/rKM9b6B2LqwSB5ToN/learning-multi-level-features-with-matryoshka-saes> [visited on 05/10/2026].
-  Chanin, David, James Wilken-Smith, Tomáš Dulka, Hardik Bhatnagar, Satvik Golechha, and Joseph Bloom [Sept. 2024]. *A Is for Absorption: Studying Feature Splitting and Absorption in Sparse Autoencoders*. URL: <https://arxiv.org/abs/2409.14507>.

References iii

-  Conmy, Arthur, Augustine N. Mavor-Parker, Aengus Lynch, Stefan Heimersheim, and Adrià Garriga-Alonso [Oct. 2023]. *Towards Automated Circuit Discovery for Mechanistic Interpretability*. URL: <http://arxiv.org/abs/2304.14997> [visited on 04/02/2026].
-  Cybenko, George [1989]. "Approximation by Superpositions of a Sigmoidal Function". In: *Mathematics of Control, Signals and Systems* 2.4.
-  Elhage, Nelson, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, Roger Grosse, Sam McCandlish, Jared Kaplan, Dario Amodei, Martin Wattenberg, and Christopher Olah [Sept. 2022]. *Toy Models of Superposition*. URL: <http://arxiv.org/abs/2209.10652> [visited on 07/17/2025].
-  Elhage, Nelson, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah [Dec. 2021]. *A Mathematical Framework for Transformer Circuits*. URL: <https://transformer-circuits.pub/2021/framework/index.html>.

References iv

-  Fillingham, Sean P., Andrew Gordon, Peter Lai, Xavier Poncini, David Quarel, and Stefan Heimersheim [Nov. 2025]. *SCALAR: Benchmarking SAE Interaction Sparsity in Toy LLMs*. URL: <http://arxiv.org/abs/2511.07572> [visited on 04/02/2026].
-  Johnson, William B. and Joram Lindenstrauss [1984]. “Extensions of Lipschitz Mappings into a Hilbert Space”. In.
-  Li, Kenneth, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg [June 2024]. *Inference-Time Intervention: Eliciting Truthful Answers from a Language Model*. URL: <http://arxiv.org/abs/2306.03341> [visited on 03/31/2026].
-  Lindsey, Jack, Wes Gurnee, Emmanuel Ameisen, Brian Chen, Adam Pearce, Nicholas L. Turner, Craig Citro, David Abrahams, Shan Carter, Basil Hosmer, Jonathan Marcus, Michael Sklar, Adly Templeton, Trenton Bricken, Callum McDougall, Hoagy Cunningham, Thomas Henighan, Adam Jermyn, Andy Jones, Andrew Persic, Zhenyi Qi, T. Ben Thompson, Sam Zimmerman, Kelley Rivoire, Thomas Conerly, Chris Olah, and Joshua Batson [Mar. 2025]. *On the Biology of a Large Language Model*. URL: <https://transformer-circuits.pub/2025/attribution-graphs/biology.html>.

References v

-  Lindsey, Jack, Adly Templeton, Jonathan Marcus, Thomas Conerly, Joshua Batson, and Christopher Olah [Oct. 2024]. *Sparse Crosscoders for Cross-Layer Features and Model Diffing*. URL: <https://transformer-circuits.pub/2024/crosscoders/index.html>.
-  McDougall, Callum [Jan. 2023]. “Induction Heads - Illustrated — LessWrong”. In: URL: <https://www.lesswrong.com/posts/TvrfY4c9eaGLeyDkE/induction-heads-illustrated> [visited on 04/02/2026].
-  Nanda, Neel [Mar. 2022]. “A Longlist of Theories of Impact for Interpretability”. In: URL: <https://www.lesswrong.com/posts/uK6sQCNMw8WKzJeCQ/a-longlist-of-theories-of-impact-for-interpretability> [visited on 03/31/2026].
-  Nanda, Neel, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt [Jan. 2023]. *Progress Measures for Grokking via Mechanistic Interpretability*. URL: <https://arxiv.org/abs/2301.05217>.

References vi

-  Nanda, Neel, Andrew Lee, and Martin Wattenberg [2023]. “Emergent Linear Representations in World Models of Self-Supervised Sequence Models”. In: *Proceedings of the 6th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*. Singapore: Association for Computational Linguistics. URL: <https://aclanthology.org/2023.blackboxnlp-1.2> [visited on 03/30/2026].
-  Olah, Chris, Nick Cammarata, Ludwig Schubert, Gabriel Goh, Michael Petrov, and Shan Carter [Mar. 2020]. “Zoom In: An Introduction to Circuits”. In: *Distill* 5.3. URL: <https://distill.pub/2020/circuits/zoom-in> [visited on 03/30/2026].
-  Panickssery, Nina, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Matt Turner [July 2024]. *Steering Llama 2 via Contrastive Activation Addition*. URL: <http://arxiv.org/abs/2312.06681> [visited on 08/12/2025].
-  Park, Kiho, Yo Joong Choe, and Victor Veitch [July 2024]. “The Linear Representation Hypothesis and the Geometry of Large Language Models”. In: *arXiv preprint arXiv:2407.18757*.
-  Power, Alethea, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra [2022]. “Grokking: Generalization Beyond Overfitting on Small Algorithmic Datasets”. In: *arXiv preprint arXiv:2201.02177*.

-  Rajamanoharan, Senthooran, Arthur Conmy, Lewis Smith, Tom Lieberum, Vikrant Varma, János Kramár, Rohin Shah, and Neel Nanda [Apr. 2024]. “Improving Dictionary Learning with Gated Sparse Autoencoders”. In.
-  Sharkey, Lee, Bilal Chughtai, Joshua Batson, Jack Lindsey, Jeff Wu, Lucius Bushnaq, Nicholas Goldowsky-Dill, Stefan Heimersheim, Alejandro Ortega, Joseph Bloom, Stella Biderman, Adria Garriga-Alonso, Arthur Conmy, Neel Nanda, Jessica Rumbelow, Martin Wattenberg, Nandi Schoots, Joseph Miller, Eric J. Michaud, Stephen Casper, Max Tegmark, William Saunders, David Bau, Eric Todd, Atticus Geiger, Mor Geva, Jesse Hoogland, Daniel Murfet, and Tom McGrath [Jan. 2025]. *Open Problems in Mechanistic Interpretability*. URL: <http://arxiv.org/abs/2501.16496> [visited on 03/30/2026].
-  SuperAnnotate [2023]. *Guide to Convolutional Neural Networks*. URL: <https://www.superannotate.com/blog/guide-to-convolutional-neural-networks>.

-  Templeton, Adly, Tom Conerly, Jonathan Marcus, Jack Clark, Yan Shlegeris, Trenton Bricken, Aidan O’Gara, Emmanuel Ameisen, Adam Jermy, Baobao Zhang, Chris Olah, and Sam McCandlish [May 2024]. *Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet*. URL: <https://transformer-circuits.pub/2024/scaling-monosemanticity/> [visited on 03/26/2026].
-  Wang, Kevin, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt [Nov. 2022]. *Interpretability in the Wild: A Circuit for Indirect Object Identification in GPT-2 Small*. URL: <https://arxiv.org/abs/2211.00593>.